

Selorm Fefeti

I build the systems people have to defend under examination.

Ten years building software for firms that operate under supervision, where the product has to satisfy a regulator, a compliance officer, and an end user who never asked for any of it. Four 0-to-1 launches across enterprise and SMB, and \$1M+ in cumulative ARR delivered at Intapp, Compliance Solutions Strategies, and Dataminr.

selormfefeti@gmail.com · [linkedin.com/in/selormfefeti](https://www.linkedin.com/in/selormfefeti) · 503.421.4381 · New York, NY

SELECTED OUTCOMES

30% of client ARR	Audit Trail for Read Events, field-level access tracking for Market Abuse Regulation compliance	Intapp
Weeks to seconds	Environment provisioning moved from an internal queue into a client-facing self-service portal	Intapp
–50% tickets	Permissioned internal tooling so client-facing teams could resolve queries without engineering	Intapp
\$600K ARR / 35 clients	Employee trade monitoring module, built 0-to-1 against SEC regulations	CSS
\$300K annual revenue	Shares Outstanding repositioned on a proprietary ML data repository, inside 12 months	CSS

CASE STUDIES

Turning the audit trail into a product

INTAPP · 2021–2026

SITUATION	Compliance officers had transactional logs they couldn't query. They could see that an event happened, but not who accessed which sensitive data or whether it fit that person's normal pattern. For firms under the Market Abuse Regulation, detection was entirely retrospective.
HARD PART	Not the tracking. Making the signal usable. Flag everything and compliance officers drown in noise, stop trusting the tool, and go back to reviewing after the fact.
THE CALLS	Ranked access anomalies by deviation from a user's normal scope and cadence, then retuned that ranking against what reviewers actually escalated. Argued for ClickHouse over a conventional database the team knew better, on WORM immutability, query speed during a live investigation, and a seven-year retention obligation.
RESULT	Field-level visibility and a defensible answer to a regulator. A condition of a major enterprise deal, and a flagship capability in the suite generating 30% of participating client ARR.

Weeks to seconds, without handing over the keys

INTAPP · 2021–2026

SITUATION	Platform administrators routed every configuration change through internal fulfillment: new environments, feature toggles, IP restrictions. Weeks of latency on changes that should have been straightforward.
FINDING	Research with security and technology officers at enterprise clients inverted the brief. The binding requirement wasn't speed or convenience. It was that the wrong person should never be able to make a high-stakes configuration change without explicit authorization.
THE CALL	I designed the role-based access model first and let the self-service interface follow from it. Deciding which actions required authorization was the real design work.
RESULT	Provisioning went from weeks to seconds and site management left the support queue. Because the portal knew which sites were live, it also surfaced hosting cost that had never been tracked.

FOUND IT	Not assigned to me. When I join a company I make a point of understanding the architecture and then walking it from a client's point of view. The Shares Outstanding repository produced a single figure from narrow sources, refreshed infrequently, and that figure is the denominator every threshold calculation depends on.
THE CASE	Senior technical and business stakeholders pushed back, reasonably: rearchitecting is expensive and nothing was visibly broken. So I stopped arguing about architecture and built the exhibit instead, setting the existing figure beside the one I calculated from additional sources and measuring both against the thresholds the number exists to test.
RESULT	A six-month rearchitecture, approved and delivered. The corrected model became the firm's proprietary ML data repository, which the Shares Outstanding product was then repositioned around, generating \$300K in annual revenue within twelve months.

WORKING WITH AI

Claude is my thinking partner for pressure-testing my own reasoning, catching gaps in requirements before they reach engineering, and synthesizing complex context fast. Four places it does real work:

PRODUCT DEFINITION

Drafting PRDs and then attacking them: what is ambiguous, what is unstated, what an engineer would otherwise have to guess. Far cheaper to find before engineering starts.

GTM STRATEGY

Segmentation, positioning, pricing narratives, and competitive teardowns, with the model taking the buyer's side so the weak parts surface early.

CUSTOMER RESEARCH

Synthesizing interviews, support threads, and call notes into patterns, then having the model argue against the pattern instead of confirming it.

MY OWN PIPELINE

I run my job search on a tool I built. It pulls postings from company job boards into a local database, dedupes them, and scores each role against my background with an LLM evaluator.

Tooling. I work with skills, connectors, MCP servers, and agents, composing them into workflows rather than prompting one thing at a time. It is also how I think about building AI products: the hard problems live at the seams between the model, the tools it can call, and the data it is allowed to reach.

HOW I WORK

I find the structural problem nobody assigned me to find, build the case for fixing it, and drive it through without waiting for someone to hand me the answer or the authority to act. What makes that work is arguing in the language of client use cases rather than opinions. A six-month rearchitecture gets approved because someone can name the client exposure it removes, not because a product manager believes in it. I also work directly with clients across financial services and legal, from pre-sales demos through onboarding and support. As subject-matter expert at industry events I generated \$500K in pipeline and 20+ qualified demos.

I work best in high-stakes environments where the output has to be trustworthy enough to act on, not just technically impressive.

ALSO SHIPPED

- **Apptruism** (own time) — charities ranked on what their public IRS filings show, close to 20,000 organizations scored, each with a confidence figure for how much of the score the data actually supports. [apptruism.streamlit.app](#)
- **LLMs in production for compliance review** (Intapp) — real-time risk detection, pattern recognition, and automated compliance checks on Azure OpenAI, cutting manual review time.
- **Form CRS filing for the SMB segment** (CSS) — product lead for the firm's first customizable filing solution; vision, roadmap, go-to-market, and pricing.
- **First industry-specific product** (Dataminr) — early-stage hire on the launch team, working across product, business, and client-facing functions.

BACKGROUND

Product 0-to-1 launches · AI/ML productization · Data analysis & management · GTM strategy · Roadmap planning · User research

Tools Azure OpenAI · SQL · Tableau · Jira · Aha! · Asana · Confluence · Figma · Grafana · Smartsheet

Education Reed College, B.A. Economics. Thesis: a Tobit model of game-day attendance across 1,786 NFL regular-season games, using point spreads and favored-team win probability as measures of outcome uncertainty.